

Table 1: LLaVA-1.5-7B — updated benchmark comparison

Snapshot: 2026-10-10 02:06:02+08:00 | Reported / reproduced rows kept separate

Method	GQA	VizWiz	SQA ^{IMG}	VQA ^{Text}	POPE-F1	MME	MMB ^{EN}	MMB ^{CN}	MMVet	Avg ₇
Full model										
FULL (reported)	61.9	55.6*	69.6	58.2	85.9	1513.4	64.7	58.1	31.3	67.7
FULL (reproduced)	61.9	—	69.5	58.2	85.9	1507.1	—	—	—	—
Retain 128 tokens										
FastV	55.3	51.9	68.7	55.6	70.0	1343.6	62.5	54.8	26.5	62.0
PDrop	56.8	49.4	69.5	56.1	77.9	1421.9	62.1	55.0	27.1	64.1
SparseVLM	57.0	49.7	69.3	56.0	83.3	1408.3	62.2	56.1	30.1	64.9
PruMerge+	58.0	53.7	68.6	54.3	83.2	1411.9	62.0	55.1	29.9	64.5
TRIM	58.7	51.6	68.3	52.2	85.5	1418.5	61.9	52.9	29.5	64.3
VisionZip	57.9	51.6	68.9	55.9	84.8	1423.7	61.5	55.9	31.3	65.2
DART	58.3	52.8	69.1	55.8	81.2	1415.1	60.7	56.0	31.1	64.6
DivPrune	59.2	57.5*	68.2	55.3	86.8	1405.1	61.5	54.8	30.6	65.2
CDPruner	59.3	57.3*	69.0	56.0	87.3	1431.4	62.4	54.7	32.1	65.8
HiPrune	57.3	55.4*	68.3	56.2	82.8	1364.4	62.2	56.4	31.2	64.5
EADP (reported)	60.0	58.0*	69.0	56.5	87.2	1439.0	62.7	55.2	32.5	66.1
EADP (reproduced)	59.7	—	69.6	56.4	87.5	1432.0	—	—	—	—
AnchorZip (Ours)	59.8	57.9*	69.2	56.6	87.5	1445.6	—	—	—	—
Retain 64 tokens										
FastV	48.0	49.1	68.7	52.1	37.2	992.1	49.6	43.3	20.1	49.8
PDrop	49.2	46.3	68.3	51.1	40.1	1004.4	48.3	37.8	19.7	49.3
SparseVLM	51.5	49.4	69.0	52.5	70.2	1192.9	58.0	49.3	25.1	58.6
PruMerge+	55.1	53.7	69.2	52.1	76.3	1310.5	59.2	51.6	28.0	61.3
TRIM	56.8	51.1	69.1	50.8	85.3	1356.1	60.5	49.7	26.5	62.9
VisionZip	56.0	52.9	69.0	55.0	77.9	1362.2	60.3	55.0	29.0	63.0
DART	55.2	53.5	68.7	54.3	74.3	1371.3	59.3	54.0	26.5	62.1
DivPrune	57.5	57.8*	68.0	54.5	85.5	1356.2	60.1	52.3	28.1	63.7
CDPruner	58.6	58.0*	68.1	54.7	87.0	1399.1	61.1	53.2	29.5	64.7
HiPrune	53.6	55.4*	68.9	54.2	73.0	1236.0	59.5	53.4	27.9	60.6
EADP (reported)	59.4	59.5*	68.9	55.0	86.5	1403.6	61.9	53.1	29.4	65.0
EADP (reproduced)	—	—	68.8	54.9	—	—	—	—	—	—
AnchorZip (Ours)	59.0	59.0*	69.1	55.0	86.5	1401.4	—	—	—	—
Retain 32 tokens										
PruMerge+	51.3	53.5	68.1	48.7	68.9	1255.2	54.6	45.2	25.5	57.1
TRIM	53.7	50.7	68.0	46.5	84.2	1259.3	57.1	39.5	22.8	58.9
VisionZip	52.2	52.4	68.7	52.0	70.6	1257.1	56.8	48.9	24.8	58.9
DART	52.6	52.5	69.1	52.2	71.1	1288.4	58.5	49.3	25.5	59.6
DivPrune	54.9	57.4*	68.6	52.4	81.5	1284.9	57.6	49.1	26.3	61.2
CDPruner	57.0	57.9*	69.5	52.1	87.1	1353.0	59.1	49.6	27.8	63.2
EADP (reported)	58.2	59.3*	69.3	52.7	86.7	1347.0	59.4	50.1	30.6	63.4
EADP (reproduced)	—	—	68.8	52.5	84.1	—	—	—	—	—
AnchorZip (Ours)	57.1	59.8*	69.5	52.5	83.1	1353.0	—	—	—	—

* VizWiz validation: EADP released scoring, author confirmation reported by user; official LOO kept for audit.
VizWiz ours: all six original alpha0.5/beta2, full-guidance, max1024 predictions, before stream-wait reruns; no per-cell maximum selection.
Public alpha0/beta1 paired controls are supplementary; their EADP scores do not fill the missing same-recipe primary baseline.
SQA: shipped CQM-A, 2017 image questions. POPE: mean F1 (2910/3000/3000). NeXT budgets are nominal.
Avg7 = (GQA + SQA-IMG + TextVQA + POPE-F1 + MME/20 + MMB-EN + MMB-CN)/7; MME is perception raw.
Local Avg7 stays blank for missing/unresolved MMB. Reported baselines are references, not local reproductions.

Table 2: LLaVA-NeXT-7B — updated benchmark comparison

Snapshot: 2026-10-10 02:06:02+08:00 | Reported / reproduced rows kept separate

Method	GQA	VizWiz	SQA ^{IMG}	VQA ^{Text}	POPE-F1	MME	MMB ^{EN}	MMB ^{CN}	MMVet	Avg ₇
Full model										
FULL (reported)	62.5	60.9*	67.6	60.3	86.8	1510.9	65.8	57.3	40.0	68.0
FULL (reproduced)	—	—	67.6	60.4	—	—	—	—	—	—
Retain 640 tokens (nominal)										
FastV	58.4	54.9	67.2	57.5	80.0	1433.2	62.4	54.1	37.1	64.5
PDrop	59.5	53.8	66.9	57.1	83.2	1456.3	63.3	54.8	35.4	65.4
SparseVLM	61.4	53.6	67.3	58.7	85.7	1464.8	64.8	57.1	35.7	66.9
PruMerge+	61.1	57.9	67.4	55.2	85.5	1476.9	64.0	56.6	33.6	66.2
TRIM	61.7	54.8	67.1	55.5	85.9	1473.5	66.0	55.9	36.9	66.5
VisionZip	61.4	57.1	67.5	58.5	86.2	1492.1	65.8	58.3	39.2	67.5
DART	61.0	57.0	68.0	59.0	85.7	1444.9	65.1	57.5	37.3	66.9
DivPrune	61.9	60.6*	67.8	57.0	86.9	1469.7	65.8	57.3	38.0	67.2
CDPruner	62.3	60.8*	67.9	58.4	87.1	1474.5	66.2	57.5	40.7	67.6
HiPrune	60.7	61.2*	68.0	48.6	85.3	1475.3	67.0	58.8	38.4	66.0
EADP (reported)	62.7	60.6*	68.0	59.2	87.4	1494.7	66.5	57.9	40.1	68.1
EADP (reproduced)	62.9	—	67.8	57.9	—	—	—	—	—	—
AnchorZip (Ours)	62.6	60.7*	68.1	57.6	87.7	1462.1	64.3	—	—	—
Retain 320 tokens (nominal)										
FastV	52.1	51.3	66.3	52.9	59.9	1134.3	55.8	46.4	21.1	55.7
PDrop	53.3	49.7	66.5	50.7	63.5	1180.7	56.9	48.9	23.7	57.0
SparseVLM	57.8	54.2	67.0	56.0	77.8	1383.5	63.3	53.4	31.9	63.5
PruMerge+	58.4	57.7	67.5	54.2	79.1	1423.7	62.6	54.8	32.3	64.0
TRIM	60.2	53.5	66.4	50.9	85.7	1445.8	62.1	52.2	33.5	64.3
VisionZip	59.3	56.2	67.3	56.4	82.9	1402.7	62.7	55.6	35.2	64.9
DART	59.1	56.8	67.1	56.3	81.4	1433.5	63.8	55.4	36.3	65.0
DivPrune	61.1	60.1*	67.5	56.2	84.7	1423.3	63.9	55.7	34.8	65.8
CDPruner	61.6	59.9*	67.7	56.4	86.8	1453.0	64.1	55.7	37.9	66.4
HiPrune	57.6	61.3*	67.3	56.5	78.9	1406.8	64.2	57.0	34.7	64.5
EADP (reported)	62.2	60.4*	67.6	57.0	86.9	1491.3	64.6	55.9	39.4	67.0
EADP (reproduced)	—	—	67.2	55.4	—	—	—	—	—	—
AnchorZip (Ours)	62.3	60.1*	67.5	56.3	85.9	1438.3	62.2	—	—	—
Retain 160 tokens (nominal)										
PruMerge+	55.6	57.2	66.7	50.9	73.7	1298.3	58.3	49.1	27.6	59.9
TRIM	57.1	52.9	65.7	47.4	84.4	1281.6	61.2	45.4	30.1	60.8
VisionZip	55.8	55.5	67.7	53.4	76.0	1319.4	59.2	51.2	31.6	61.3
DART	56.3	56.7	67.5	53.6	77.2	1332.8	61.5	53.3	32.0	62.3
DivPrune	60.2	60.7*	67.1	53.7	80.0	1376.3	62.3	53.4	32.4	63.6
CDPruner	60.6	59.7*	67.5	53.2	86.3	1402.3	63.6	53.8	35.0	65.0
HiPrune	53.7	59.5*	68.7	48.6	67.7	1200.7	59.8	50.7	29.2	58.5
EADP (reported)	61.6	60.2*	67.4	54.2	86.0	1419.5	63.4	54.6	35.6	65.5
EADP (reproduced)	—	—	67.5	52.1	—	—	—	—	—	—
AnchorZip (Ours)	61.2	59.7*	67.1	53.3	82.8	1373.1	60.8	—	—	—

* VizWiz validation: EADP released scoring, author confirmation reported by user; official LOO kept for audit.
VizWiz ours: all six original alpha0.5/beta2, full-guidance, max1024 predictions, before stream-wait reruns; no per-cell maximum selection.
Public alpha0/beta1 paired controls are supplementary; their EADP scores do not fill the missing same-recipe primary baseline.
SQA: shipped CQM-A, 2017 image questions. POPE: mean F1 (2910/3000/3000). NeXT budgets are nominal.
Avg7 = (GQA + SQA-IMG + TextVQA + POPE-F1 + MME/20 + MMB-EN + MMB-CN)/7; MME is perception raw.
Local Avg7 stays blank for missing/unresolved MMB. Reported baselines are references, not local reproductions.

Table 3: Qwen3-VL-8B — updated benchmark comparison

Snapshot: 2026-10-10 02:06:02+08:00 | Reported / reproduced rows kept separate

Method	VQA ^{Text}	ChartQA	AI2D	OCRBench	HallBench	MME	MMB ^{EN}	MMB ^{CN}	DocVQA	InfoVQA	Avg ₁₀
<i>Full model</i>											
FULL (reported)	83.8	82.3	86.3	855	51.2	2440.2	86.3	84.6	94.5	66.4	84.3
<i>Retain 512 tokens</i>											
DivPrune	75.0	59.8	78.6	658	41.0	2217.5	82.1	79.4	74.6	46.0	71.3
CDPruner	74.5	62.0	76.9	660	41.6	2199.8	81.0	79.3	72.3	47.3	71.1
EADP (reported)	75.2	65.9	77.7	673	42.5	2261.1	81.6	79.7	77.7	49.9	73.1
<i>Retain 256 tokens</i>											
DivPrune	69.5	47.5	75.2	595	38.7	2185.3	80.0	78.8	55.8	38.3	65.3
CDPruner	68.5	51.9	74.1	601	41.3	2200.6	79.8	80.0	53.0	37.6	65.6
EADP (reported)	71.4	55.5	75.7	625	42.8	2202.4	80.5	78.9	62.8	42.3	68.2
EADP (reproduced)	71.0	54.8	74.5	623	41.0	2212.7	79.9	78.9	61.1	42.4	67.7
AnchorZip (Ours)	72.0	55.0	74.8	632	41.0	2202.9	79.5	78.4	63.1	43.0	68.0
<i>Retain 128 tokens</i>											
DivPrune	62.6	37.0	72.4	499	36.7	2136.3	77.0	76.6	39.5	33.8	59.2
CDPruner	60.6	40.4	71.7	515	34.7	2153.8	77.9	77.5	37.8	32.5	59.2
EADP (reported)	66.1	43.6	73.6	550	40.2	2200.2	79.9	78.4	45.3	35.3	62.7
AnchorZip (Ours)	67.4	42.6	73.1	555	38.4	2148.5	79.0	77.2	46.0	36.0	62.3
<i>Retain 64 tokens</i>											
AnchorZip (Ours)	59.2	33.0	71.0	495	34.7	2088.3	78.0	77.2	32.5	30.5	57.0

† Old complete result; repaired rerun pending. — Missing or unresolved protocol. * VizWiz validation.
SQA: shipped CQM-A, 2017 image questions. POPE: mean F1 (2910 / 3000 / 3000). NeXT budgets are nominal.
LLaVA MME: perception; Qwen MME: perception + cognition. Reported baselines are references, not our reproductions.
Avg10 = (8 percentage columns + OCRBench/10 + MME/20)/10. Qwen rows are prior completed legacy runs.
EADP K256 HallBench: scoring audited; generation-bank provenance unresolved. No local EADP K128/K512 comparison.